Final Project
STAT 4110 Statistical Simulation
Winter 2026 Semester

School of Mathematical and Computational Sciences
University of Prince Edward Island
Ali Raisolsadat

January 10, 2026

Overview

This project is designed to familiarize students with real-world applications of statistical simulation, stochastic modelling, and analytical reasoning. The available options for the final project may involve developing a computational model, implementing simulation techniques covered in class, and using simulation-based or analytical methods to study system behaviour. Students are free to use any dataset or programming language of their choice, unless otherwise specified.

Choose **one** of the following three project directions:

1. Use your modelling skills to model and implement **one of the following three** real world problems:

   - Forest Fire Simulation – Poisson and Markov Chain Simulation
   - Numerical Option Pricing: Crank–Nicolson PDE vs. Monte Carlo Simulation
   - Bayesian Neural Networks using Markov Chain Monte Carlo
   - Your very own problem and modelling. Some examples are: Queueing Problems, Molecule Collisions in a Reactor, Prey/Predator using Poisson Point Estimates, 1D and 2D Ising Model.

2. Replicating the full design and implementation of **one** published paper related to our course material

3. Literature review of **5–15 papers** on a topic related to the course

Projects will be completed in teams of three members. Each team will be responsible for defining the scope, data sources, and motivation of their work.

Deliverables

For any of the project choices, the deliverables are as follows:

1. **Introductory Summary** (2 pages max, 5/30 points): Includes a brief description of the chosen project, references and papers that the group will use, and the roles of each student in the group.

2. **Summary Report** (5 pages max, 20/30 points): A detailed report summarizing the approach, implementation, results, and discussion. Include any supporting figures, tables, or appendices as needed. You are also responsible to attach a ZIP file containing all code used in the project. If you are using Git (GitHub or GitLab), please add me as the owner, so I can review your code.

3. **Anonymous Peer Review** (5/30 Points): A peer review questionnaire with 5 grading questions and one paragraph question describing your peers contribution to the project.

   Disclaimer: The use of LLMs, including GPT, is permitted for coding assistance only. Any use of such tools must be explicitly acknowledged in the report.

Project Choice 1: Modelling Option

Purpose

The purpose of this project choice is to give you hands-on experience in developing and analysing statistical/mathematical models for real-world problems. By working on a modelling-based project, you will learn how to formalize assumptions, implement computational methods, and interpret model outcomes. This choice emphasizes on using the material learnt throughout the course and understanding the behaviour of complex systems under uncertainty

Learning Objectives

- Design and implement a computational model for a selected application.

- Apply simulation or numerical methods to approximate solutions when analytic solutions are impossible.

- Compare model predictions with alternative approaches or benchmarks.

- Communicate methodology, results, and insights effectively through written reports.

Scope

You are expected to choose **one of the four** modelling-based problems:

- Forest Fire Simulation – Poisson and Markov Chain Simulation

- Numerical Option Pricing: Crank–Nicolson PDE vs. Monte Carlo Simulation

- Bayesian Neural Networks using Markov Chain Monte Carlo

- Your very own problem and modelling. Some examples are: Queueing Problems, Molecule Collisions in a Reactor, Prey/Predator using Poisson Point Estimates, 1D and 2D Ising Model.

As a team, you are responsible for selecting appropriate datasets, defining the project scope, and documenting their methodology and results.

Option 1: Forest Fire Simulation

Forest fires are a significant environmental hazard that can have devastating ecological and economic impacts. Fires can ignite naturally, such as through lightning strikes, or by human activity. The propagation of a fire through a forest depends on multiple factors including tree density, spatial arrangement of trees, terrain, and climatic conditions. Understanding these dynamics through simplified mathematical models allows forest managers to identify high-risk areas, estimate the time-to-burn, and develop effective mitigation strategies.

Spatial heterogeneity in forests plays a crucial role in fire behaviour. Trees are rarely uniformly distributed; regions with fertile soil or favourable growth conditions often have higher tree densities, forming natural clusters that are more susceptible to rapid fire spread. Stochastic models that incorporate both the spatial distribution of trees and probabilistic fire spread can provide quantitative insights into these dynamics.

In this project, you are asked to construct a statistical model for probabilistic forest fire spread in a 2D square forest of size 100×100 km^2. The goal is to quantify fire propagation probabilities, expected number of burnt trees, and the time-to-burn under varying conditions.

Objectives

This problem emphasizes both mathematical modelling and probabilistic reasoning. Students are expected to:

1. Model a spatially heterogeneous forest using inhomogeneous Poisson point process (IPP).

2. Represent fire spread using a discrete-time Markov chain over trees as states.

3. Perform Monte Carlo simulations to estimate expected fire outcomes, such as:

   - Probability that a tree burns given another tree ignites
   - Expected number of burnt trees
   - Expected time-to-burn for the forest

4. Visualize both tree density and dominant fire pathways.

5. Analyse spatial patterns of fire risk and identify areas with consistently high risk.

Step 1: Forest Model – Inhomogeneous Poisson Point Process

Let the forest domain be a square:

$$D = [0, 100] \times [0, 100]$$

Define a radially symmetric Gaussian intensity function:

$$\lambda(x, y) = \lambda_{\max} \exp\left(-\frac{(x - c)^2 + (y - c)^2}{2\sigma^2} \right)$$

where $c = 50$ is the center of the forest and σ controls the concentration of trees near the center.

Tree Sampling Algorithm:

1. Generate a Poisson number of candidate trees:

$$N_{\text{prop}} \sim \text{Poisson}(\lambda_{\max} \cdot |D|)$$

2. Assign candidate coordinates uniformly in D.

3. Accept each candidate with probability

$$p_{\text{accept}} = \frac{\lambda(x_i, y_i)}{\lambda_{\max}}$$

4. Continue until exactly N_{trees} are accepted.

This ensures that tree locations reflect the desired spatial heterogeneity.

Step 2: Fire Spread – Markov Chain Dynamics

Define the state of each tree i at timestep t as

$$S_i(t) \in \{0, 1\}, \quad 0 = \text{unburnt}, \ 1 = \text{burning}$$

Let $\mathcal{N}_i = \{j : \text{distance}(i, j) \le r_{\text{fire}}\}$ denote the neighbors of tree i within a specified fire radius.
Transition probability: If tree i is burning at time t, each neighbor $j \in \mathcal{N}_i$ ignites at time $t + 1$ with probability
$$\Pr(S_j(t+1) = 1 \mid S_j(t) = 0, S_i(t) = 1) = p_{\text{fire}}$$

- Trees remain burning once ignited.

- Fire propagates in discrete timesteps until no new trees ignite or a maximum number of timesteps is reached.

Efficient Neighbor Search: In practical implementation, computing $\mathcal{N}_i$ for every tree by brute-force distance checks is $\mathcal{O}(N^2)$, which becomes costly for large N. Python's `scipy.spatial.cKDTree` allows efficient spatial queries:

- `cKDTree(tree_pts)` constructs a spatial index of all tree coordinates in $O(N \log N)$ time.

- `query_ball_tree(tree, r)` returns a list of neighbors for each tree within radius r_{fire}, effectively computing $\mathcal{N}_i$ for all i.

- Mathematically, this corresponds to finding all j such that

$$||\mathbf{x}_i - \mathbf{x}_j||_2 \le r_{\text{fire}}.$$

Step 3: Monte Carlo Estimation of Fire Transition Probabilities

To quantify fire dynamics probabilistically:

1. Run N_{MC} independent fire simulations.

2. Record all directed transitions $(i \to j)$, representing tree i igniting tree j.

3. Compute the empirical Markov transition matrix

$$\hat{P}_{ij} = \frac{\text{count of transitions } i \to j}{N_{\text{MC}}}$$

4. The matrix $\hat{P}$ encodes the probability that tree j burns given tree i burns.

Step 4: Visualization and Analysis

- Represent each tree as a node in 2D coordinates.

- Draw directed edges $(i \to j)$ with thickness proportional to $\hat{P}_{ij}$ for dominant transitions.

- Colour-code nodes according to the cumulative influence or total probability of being burnt.

- Use this visualization to identify spatial fire risk patterns and likely fire propagation pathways.

For large forests, visualizing all pairwise transitions with arrows can be cluttered. Instead, an aggregated heatmap can be used:

- Define a 2D grid over the forest domain (mesh of cells).

- Compute **outgoing fire influence** for each tree as

$$f_i = \sum_j \hat{P}_{ij},$$

  representing the expected number of neighbors ignited by tree i.

- Aggregate these influences into each grid cell by averaging the f_i of trees falling within that cell.

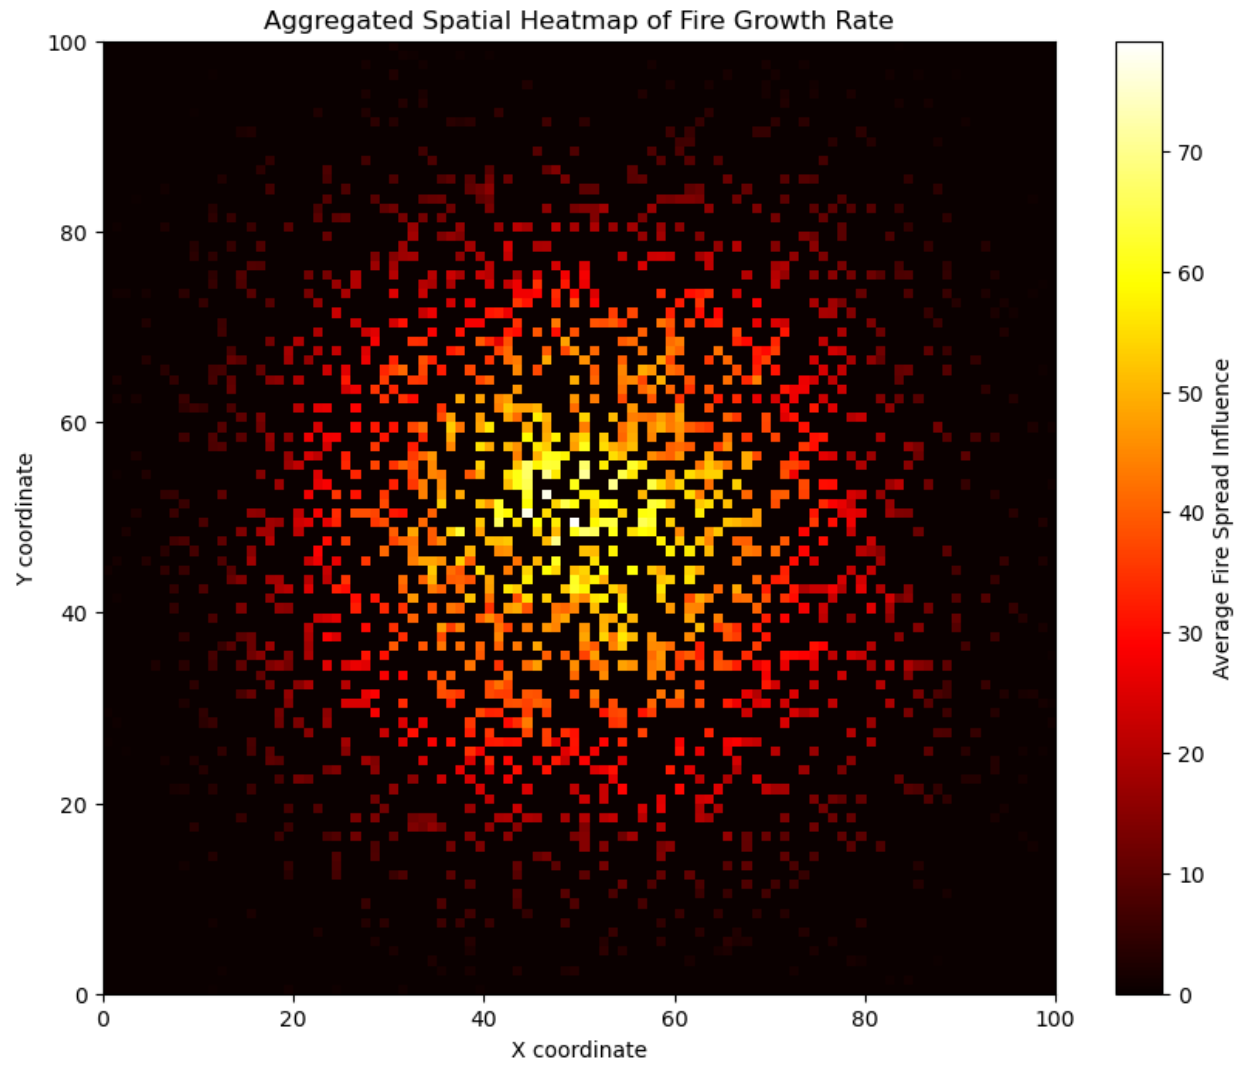**Step 5: Analysis of Results**

Students should analyze:

- Which trees consistently act as ignition points.

- Areas of high fire connectivity and risk.

- How spatial heterogeneity affects the spread of fire.

- Potential interventions to mitigate fire risk (e.g., thinning high-density clusters).

Suggested Parameters

Parameter	Value	Notes
Forest size	100 km	2D square domain
Maximum tree intensity ($\lambda_{\max}$)	0.05	High density near center
Fire radius (r_{fire})	20 km	Neighbor search radius
Maximum timesteps	120	Time resolution per step
Monte Carlo trials (N_{MC})	500–2000	Per fire scenario
Random seed	42	For reproducibility

Directed Fire Spread
(Dominant Markov Transitions)

Aggregated Spatial Heatmap of Fire Growth Rate

Option 2: Black–Scholes Model, Monte Carlo Pricing, and Crank–Nicolson Finite Differences

1. What is the Black–Scholes model and why use it?

The Black–Scholes model is a continuous-time stochastic model for the evolution of a financial asset price S_t. It assumes that the asset price follows a geometric Brownian motion (GBM):

$$dS_t = \mu S_t \, dt + \sigma S_t \, dW_t$$

where

- μ is the instantaneous expected return (drift),

- $\sigma > 0$ is the volatility (instantaneous standard deviation),

- W_t is a standard Brownian motion.

The model is used because (i) it is analytically tractable in many cases, (ii) under risk-neutral valuation it leads to a linear parabolic PDE for derivative prices, and (iii) it captures multiplicative stochastic growth and log-normal marginal distributions for S_t. The Black–Scholes framework underpins much of modern quantitative finance and provides a baseline for comparing numerical methods (PDE solvers, Monte Carlo, lattice models, etc.).

2. Risk-neutral pricing and derivation of the Black–Scholes PDE

Consider a derivative security with payoff at maturity T given by $\Phi(S_T)$. The *no-arbitrage* (risk-neutral) pricing formula states that the arbitrage-free price at time t is the discounted risk-neutral expectation:

$$V(S_t, t) = e^{-r(T-t)} \mathbb{E}^{\mathbb{Q}}\big[\, \Phi(S_T) \mid S_t \, \big]$$

where r is the risk-free rate and $\mathbb{Q}$ denotes the risk-neutral measure. Under the risk-neutral measure the asset dynamics become

$$dS_t = r S_t \, dt + \sigma S_t \, dW_t^{\mathbb{Q}}$$

Applying Itô's formula to $V(S_t, t)$ gives

$$dV = \left(\frac{\partial V}{\partial t} + rS\frac{\partial V}{\partial S} + \tfrac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} \right) dt + \sigma S \frac{\partial V}{\partial S} dW_t^{\mathbb{Q}}$$

Under risk-neutral valuation, the discounted derivative price must be a martingale, which leads to the Black–Scholes backward PDE:

$$\boxed{\frac{\partial V}{\partial t} + \tfrac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + rS\frac{\partial V}{\partial S} - rV = 0}$$

Terminal condition: $V(S, T) = \Phi(S)$.

3. Probabilistic (Monte Carlo) representation

The derivative price can be approximated by the risk-neutral expectation:

$$V(S_0, 0) = e^{-rT} \mathbb{E}^{\mathbb{Q}}[\Phi(S_T) \mid S_0]$$

Under GBM,

$$S_T = S_0 \exp\left(\big(r - \tfrac{1}{2}\sigma^2\big)T + \sigma\sqrt{T}\, Z \right), \qquad Z \sim \mathcal{N}(0, 1)$$

3.1 Euler–Maruyama discretization

Discretize the interval $[0, T]$ into N steps $\Delta t = T/N$. The Euler–Maruyama scheme:

$$S_{n+1} = S_n + rS_n\Delta t + \sigma S_n\sqrt{\Delta t}Z_n \quad Z_n \sim N(0,1)$$

or the log-Euler form:

$$S_{n+1} = S_n \exp\left((r - \tfrac{1}{2}\sigma^2)\Delta t + \sigma\sqrt{\Delta t}\,Z_n\right)$$

3.2 Monte Carlo pricing of European options

Simulate M paths $S_T^{(m)}$ and compute:

$$V_0 \approx e^{-rT}\frac{1}{M}\sum_{m=1}^{M}\Phi\big(S_T^{(m)}\big)$$

4. Numerical solution of the PDE: Crank–Nicolson finite differences

Finite-difference methods discretize the PDE on a grid in asset price S and time t. Crank–Nicolson (CN) is a widely used implicit scheme that is second-order accurate in both space and time and is unconditionally stable (for the linear parabolic PDEs we consider).

Below we derive the CN discretization and obtain the tridiagonal linear system that must be solved at each time step.

4.1 Domain and grid

Truncate the infinite spatial domain $S \in (0, \infty)$ to a large finite domain $S \in [S_{\min}, S_{\max}]$. A common choice is $S_{\min} = 0$ and $S_{\max} = LS_0$ with $L = 3$ or 4, chosen so the option value is insensitive to further increases in $S_{\max}$.

Define a uniform spatial grid with $M + 1$ nodes:

$$S_i = S_{\min} + i\Delta S \qquad \Delta S = \frac{S_{\max} - S_{\min}}{M} \quad i = 0, 1, \ldots, M$$

Define a temporal grid stepping backward from T to 0 with N time-steps:

$$t^n = n\Delta t \qquad \Delta t = \frac{T}{N} \quad n = 0, 1, \ldots, N$$

We will denote the numerical approximation to $V(S_i, t^n)$ by V_i^n. (Because the PDE is typically integrated backward in time from $t = T$ to $t = 0$, many authors index time reversely; here we take $n = 0$ for time $t = 0$ and note the CN update moves from known V^n to V^{n+1}.)

4.2 Spatial derivatives approximation

Approximate first and second derivatives by centered finite differences:

$$\left.\frac{\partial V}{\partial S}\right|_{S_i t^n} \approx \frac{V_{i+1}^n - V_{i-1}^n}{2\Delta S} \qquad \left.\frac{\partial^2 V}{\partial S^2}\right|_{S_i t^n} \approx \frac{V_{i+1}^n - 2V_i^n + V_{i-1}^n}{(\Delta S)^2}$$

4.3 Semi-discrete operator

Plugging these into the PDE, at grid node (S_i, t^n) the differential operator

$$\mathcal{L}V := \tfrac{1}{2}\sigma^2 S^2 V_{SS} + rSV_S - rV$$

is approximated by

$$(\mathcal{L}V)_i^n \approx \tfrac{1}{2}\sigma^2 S_i^2 \frac{V_{i+1}^n - 2V_i^n + V_{i-1}^n}{(\Delta S)^2} + rS_i \frac{V_{i+1}^n - V_{i-1}^n}{2\Delta S} - rV_i^n$$

It is convenient to rewrite this as a linear combination

$$(\mathcal{L}V)_i^n \approx \alpha_i^n V_{i-1}^n + \beta_i^n V_i^n + \gamma_i^n V_{i+1}^n$$

with coefficients (dropping the explicit time index for the coefficients since they depend only on S_i):

$$\alpha_i := \frac{1}{2}\sigma^2 S_i^2 \frac{1}{(\Delta S)^2} - \frac{rS_i}{2\Delta S}$$

$$\beta_i := -\frac{1}{2}\sigma^2 S_i^2 \frac{2}{(\Delta S)^2} - r$$

$$\gamma_i := \frac{1}{2}\sigma^2 S_i^2 \frac{1}{(\Delta S)^2} + \frac{rS_i}{2\Delta S}$$

Thus

$$(\mathcal{L}V)_i^n \approx \alpha_i V_{i-1}^n + \beta_i V_i^n + \gamma_i V_{i+1}^n$$

4.4 Crank–Nicolson time discretization

Crank–Nicolson averages the spatial operator at times t^n and t^{n+1}. The time derivative is approximated by

$$\frac{V_i^{n+1} - V_i^n}{\Delta t} \approx \frac{1}{2}\big((\mathcal{L}V)_i^{n+1} + (\mathcal{L}V)_i^n\big).$$

Rearrange to obtain the CN update equation:

$$V_i^{n+1} - \frac{\Delta t}{2}(\mathcal{L}V)_i^{n+1} = V_i^n + \frac{\Delta t}{2}(\mathcal{L}V)_i^n$$

Substituting the finite-difference representation of $\mathcal{L}$ yields a linear relation for interior nodes $i = 1, \ldots, M-1$:

$$V_i^{n+1} - \frac{\Delta t}{2}\big(\alpha_i V_{i-1}^{n+1} + \beta_i V_i^{n+1} + \gamma_i V_{i+1}^{n+1}\big) = V_i^n + \frac{\Delta t}{2}\big(\alpha_i V_{i-1}^n + \beta_i V_i^n + \gamma_i V_{i+1}^n\big)$$

Collect terms for V_{i-1}^{n+1}, V_i^{n+1}, V_{i+1}^{n+1} on the left-hand side and known V^n terms on the right-hand side. Define coefficients:

$$A_i := -\frac{\Delta t}{2}\alpha_i$$

$$B_i := 1 - \frac{\Delta t}{2}\beta_i \qquad \text{(left-hand side coefficients)}$$

$$C_i := -\frac{\Delta t}{2}\gamma_i$$

and

$$\widetilde{A}_i := \frac{\Delta t}{2}\alpha_i$$

$$\widetilde{B}_i := 1 + \frac{\Delta t}{2}\beta_i \qquad \text{(right-hand side coefficients)}$$

$$\widetilde{C}_i := \frac{\Delta t}{2}\gamma_i$$

Then the CN equation becomes, for each interior i,

$$A_i V_{i-1}^{n+1} + B_i V_i^{n+1} + C_i V_{i+1}^{n+1} = \widetilde{A}_i V_{i-1}^n + \widetilde{B}_i V_i^n + \widetilde{C}_i V_{i+1}^n$$

4.5 Matrix form

Stack the interior unknowns into a vector

$$\mathbf{V}^n = \begin{bmatrix} V_1^n \\ V_2^n \\ \vdots \\ V_{M-1}^n \end{bmatrix}$$

For each time-step the system can be written compactly as

$$\boxed{A\,\mathbf{V}^{n+1} \;=\; B\,\mathbf{V}^n \;+\; \mathbf{d}^n}$$

where A and B are $(M-1) \times (M-1)$ tridiagonal matrices with entries given by the A_i, B_i, C_i and $\widetilde{A}_i, \widetilde{B}_i, \widetilde{C}_i$ coefficients, and $\mathbf{d}^n$ is a vector that collects boundary contributions. Explicitly, A has the form

$$A = \begin{bmatrix} B_1 & C_1 & 0 & \cdots & & 0 \\ A_2 & B_2 & C_2 & \ddots & & \vdots \\ 0 & \ddots & \ddots & \ddots & & 0 \\ \vdots & \ddots & A_{M-2} & B_{M-2} & C_{M-2} \\ 0 & \cdots & 0 & A_{M-1} & B_{M-1} \end{bmatrix}$$

and B is similarly

$$B = \begin{bmatrix} \widetilde{B}_1 & \widetilde{C}_1 & 0 & \cdots & & 0 \\ \widetilde{A}_2 & \widetilde{B}_2 & \widetilde{C}_2 & \ddots & & \vdots \\ 0 & \ddots & \ddots & \ddots & & 0 \\ \vdots & \ddots & \widetilde{A}_{M-2} & \widetilde{B}_{M-2} & \widetilde{C}_{M-2} \\ 0 & \cdots & 0 & \widetilde{A}_{M-1} & \widetilde{B}_{M-1} \end{bmatrix}$$

The boundary vector $\mathbf{d}^n$ accounts for known values V_0^n and V_M^n appearing in the finite-difference stencils for the first and last interior nodes. For example, the first equation (for $i = 1$) includes terms proportional to V_0^n and V_0^{n+1}; move these to the right-hand side to obtain the appropriate entry in $\mathbf{d}^n$.

Boundary conditions. Reasonable boundary conditions for European call options:

$$V(S_{\min} = 0, t) = 0 \qquad \text{(call is worthless at zero)}$$
$$V(S_{\max}, t) \approx S_{\max} - Ke^{-r(T-t)}$$

Use these values to compute contributions to $\mathbf{d}^n$. If one uses $S_{\min} = 0$ then $V_0^n = 0$ simplifies the first equation. For the last equation, known V_M^n provides a known term on the right-hand side.

5. Implementation notes and practical considerations

Defining the option and model. Choose a standard option (e.g., European call or put) and clearly define all parameters (S_0, K, r, σ, T, dividends). For European options, the analytical Black–Scholes solution can serve as a reference.

Accuracy comparison.

- **Monte Carlo (MC):** Simulate N paths of the underlying asset, compute discounted payoff averages, and measure absolute or relative error against the analytical solution. Optionally, report standard error of the estimator.

- **Crank–Nicolson (CN):** Construct a time–price grid (Δt, ΔS) and solve the PDE. Compute the error relative to the analytical or highly accurate numerical solution.

Compare errors for both methods and observe convergence under grid refinement (CN) or increasing number of paths (MC).

Efficiency comparison.

- **Monte Carlo:** Runtime scales linearly with the number of simulated paths. Error decreases as $1/\sqrt{N}$. Record CPU time for each N.

- **Crank–Nicolson:** Runtime scales with $M \times N$, where M is the number of time steps and N the number of spatial grid points. Finer grids improve accuracy faster than MC for the same runtime in low-dimensional problems.

Generate work–precision plots (error vs CPU time) to visualize efficiency trade-offs.

Practical considerations.

- MC is more suitable for high-dimensional or path-dependent options; variance reduction techniques improve efficiency.

6. Tasks

You are asked to implement both Monte Carlo and Crank–Nicolson methods for European call options and perform a comparative study.

1. **Constants table (suggested values):**

Parameter	Value
Initial asset price S_0	100
Strike K	110 (high K values for variance reduction)
Time to maturity T	1 year
Volatility σ	0.2, 0.4
Risk-free rate r	0.01, 0.05
Number of time steps N	100
Number of Monte Carlo paths M	50,000

2. **Experiments:**

- Compute option prices for all combinations of σ and r.

- Compare Monte Carlo estimated prices with Crank–Nicolson numerical PDE prices.

- Measure CPU time and compare efficiency of Monte Carlo vs CN.

- Use variance reduction techniques for the MC estimator. Compare efficiency with new estimators.

- Plot option price vs volatility, option price vs interest rate.

Option 3: Bayesian Neural Networks via MCMC

1. Background and Motivation

Traditional neural networks optimize weights to single point estimates, usually via gradient descent. This approach does not quantify uncertainty, which is important when:

- making decisions in risk-sensitive applications (medicine, finance, engineering)

- detecting anomalies or rare events

- forecasting with small datasets

- requiring confidence intervals or credible predictions

Bayesian neural networks (BNNs) treat the network weights as random variables with prior distributions. The posterior distribution of the weights given observed data allows us to capture uncertainty:

$$p(w \mid X, y) = \frac{p(y \mid X, w)\, p(w)}{p(X, y)},$$

where w are the weights, X the inputs, y the observed outputs, $p(w)$ the prior, and $p(y \mid X, w)$ the likelihood.

Because the posterior is usually intractable, we approximate it using Markov Chain Monte Carlo (MCMC) methods such as:

- Metropolis–Hastings

- Gibbs sampling

- Hamiltonian Monte Carlo (HMC)

2. Model Structure

Consider a simple feedforward neural network with input x, one hidden layer of H neurons, and output y. Denote the network function as $f(x; w)$.

Priors: Assign independent Gaussian priors to weights and biases:

$$w \sim \mathcal{N}(0, \sigma_w^2), \quad b \sim \mathcal{N}(0, \sigma_b^2).$$

Likelihood (for regression): Assume Gaussian noise on outputs:

$$y \mid x, w \sim \mathcal{N}(f(x; w), \sigma^2).$$

Posterior: The posterior distribution of the weights is proportional to:

$$p(w \mid X, y) \propto p(y \mid X, w)p(w).$$

Posterior predictive distribution: For a new input x_{new}, approximate the predictive distribution by averaging over posterior samples:

$$\hat{y}(x_{\text{new}}) \approx \frac{1}{S} \sum_{s=1}^{S} f(x_{\text{new}}; w^{(s)}),$$

where $w^{(1)}, \ldots, w^{(S)}$ are sampled from the posterior using MCMC.

3. Datasets

You can use any dataset, but it is recommended to choose **small or simple datasets** so that MCMC sampling completes in reasonable time. Examples include:

- **Tabular / regression:** Boston Housing, Diabetes, or synthetic datasets such as $y = \sin(x) + \epsilon$.

- **Time series:** Stock prices (subset), daily sales, or temperature data.

- **Classification / small images:** Iris dataset, or a subset of MNIST (e.g., 1000–5000 images).

The key is to keep the dataset size manageable for a few thousand MCMC iterations.

4. Suggested Experiments

- Vary the prior variance σ_w^2 and observe the effect on predictive uncertainty.

- Change the noise variance σ^2 and compare predictive intervals.

- Compare performance and uncertainty estimates between MCMC BNN and standard NN.

- Optional: Vary the number of hidden neurons or layers and observe convergence of MCMC samples.

5. Tasks

1. Implement a Bayesian neural network using MCMC.

2. Generate posterior samples for the weights and biases.

3. Compute posterior predictive distributions for test inputs.

4. Investigate how predictions and uncertainty change when:

    - the prior variance changes
    - the output noise changes
    - the network size changes

5. Compare accuracy and efficiency with a standard neural network.

Option 2: Replicating the Design and Implementation of a Published Paper

Background and Motivation

Replication is a core aspect of scientific research. By attempting to reproduce the results of a published study, you gain a deeper understanding of modelling assumptions, computational methods, and analysis techniques. This option allows you to connect course concepts to real-world research and critically evaluate methodological choices.

Project Task

You will:

- Select a published paper related to stochastic modelling, Monte Carlo and/or Monte Carlo simulation or other topics covered in the course.

- Carefully review the paper, focusing on the problem statement, assumptions, model structure, and computational approach.

- Implement the methods or experiments described in the paper using a programming language of their choice.

- Compare your results with the published results and analyse any discrepancies.

- Document challenges encountered, adaptations made, and lessons learned from the replication process.

Learning Objectives

- Understand the translation of theoretical models into practical computational methods.

- Critically evaluate published results and methodologies.

- Develop practical skills in programming, numerical methods, and model validation.

- Gain experience in scientific documentation and reproducibility.

Option 3: Literature Review of 5–15 Papers

Background and Motivation

A literature review allows you to synthesize knowledge on a particular topic, identify trends, gaps, and open questions, and develop critical evaluation skills. This project option emphasizes scholarly analysis rather than implementation, and is ideal for you interested in exploring broader research directions.

Project Task

You will:

- Select a coherent topic or research question related to the course (e.g., stochastic simulations, Bayesian inference, neural network modelling, queueing systems).

- Search for and read 5–15 relevant academic papers, preprints, or technical reports.

- Summarize key methods, results, assumptions, and conclusions of each paper.

- Compare approaches, highlight similarities and differences, and identify gaps in the literature.

- Discuss potential future directions or applications suggested by the literature.

Learning Objectives

- Develop skills in critical reading, synthesis, and scholarly writing.

- Identify research trends and gaps within a focused area.

- Learn to communicate complex ideas clearly in written and oral form.

- Gain familiarity with academic databases, citation practices, and literature evaluation techniques.