

Chapter 1 - Random Number Generation

Generating Random Variables in R and Python. Discrete Distributions.

Prof. Alex Alvarez, Ali Raisolsadat

School of Mathematical and Computational Sciences
University of Prince Edward Island

Randomness has played an important role for centuries in games, gambling, and statistics.

Historically, random numbers were generated manually or mechanically (e.g., spinning wheels, rolling dice, shuffling cards, or using random number tables).

The main building block for all the simulation algorithms that we will study in this course is the generation of random numbers with a given distribution. This includes both discrete and continuous random variables.

If we are using a powerful statistical software such as **R** (but also **Python** or **MATLAB**, and most used programming languages) there are already some built-in functions that can help us generate random numbers with well known probability distributions.

On the other hand, we will also learn some techniques to deal with arbitrary probability distributions (generally not supported by any software).

The algorithms used by most software packages to generate random numbers with a given distribution are **deterministic** algorithms, meaning that the generated random numbers are not “truly random”. We call them **pseudo random number generators** (PRNG).

In other words, a PRNG is an algorithm which outputs a sequence of random numbers with properties that are very similar to the properties of sequences of true random numbers. PRNGs are fast and efficient, and for most applications they generate sequences of random numbers that can replace true random numbers without major issues.

Example: The most common method generates numbers starting from an initial value x_0 (the **seed**) and recursively computes successive values x_n for $n \geq 1$ using

$$x_n = ax_{n-1} \mod m$$

where a and m are positive integers. The remainder operation ensures that $x_n \in \{0, 1, \dots, m-1\}$. On a 32-bit machine (with the first bit as a sign bit), the choice $m = 2^{31} - 1$ and $a = 7^5 = 16807$ produces desirable properties.

Some Probability Distributions Supported by R and Python

Distribution	Name in R	Name in Python (SciPy)
Uniform distribution	<i>unif</i>	<i>uniform</i>
Normal distribution	<i>norm</i>	<i>norm</i>
Poisson distribution	<i>pois</i>	<i>poisson</i>
Binomial distribution	<i>binom</i>	<i>binom</i>

In **R**, the names of functions associated with these probability distributions are constructed by adding a prefix letter to the distribution name:

- **r**: random number generation (e.g., *runif*)
- **d**: density (or probability mass function for discrete distributions)
- **p**: cumulative distribution function (CDF)
- **q**: quantile function

Some Probability Distributions Supported by R and Python

Distribution	Name in R	Name in Python (SciPy)
Uniform distribution	<i>unif</i>	<i>uniform</i>
Normal distribution	<i>norm</i>	<i>norm</i>
Poisson distribution	<i>pois</i>	<i>poisson</i>
Binomial distribution	<i>binom</i>	<i>binom</i>

In **Python (SciPy and Numpy)**, these distributions are typically accessed via *scipy.stats*. For example:

- *stats.uniform.rvs(size=10)* generates random samples.
- *stats.uniform.pdf(x)* evaluates the probability density function.
- *stats.uniform.cdf(x)* computes the CDF.
- *stats.uniform.ppf(q)* gives the quantile function (percent-point function).

Example: *runif(10)* in **R** or *stats.uniform.rvs(size=10)* in **Python** generate 10 random numbers from a uniform distribution.

Example: Generate a random number from a random variable X with discrete probability distribution given by the p.m.f. $p(i) = i/6$ for $i = 1, 2, 3$. That is $P(X = 1) = 1/6$, $P(X = 2) = 1/3$ and $P(X = 3) = 1/2$.

Algorithm

1. Generate $U \sim U(0, 1)$
2. Define $X = \begin{cases} 1 & \text{if } 0 \leq U < 1/6 \\ 2 & \text{if } 1/6 \leq U < 1/2 \\ 3 & \text{if } 1/2 \leq U \leq 1 \end{cases}$
3. Return X

Generation of Random Numbers with Discrete Distributions

Code

R

```
1  U <- runif(1, min=0, max=1)
2  if (U < 1/6) {
3      X <- 1
4  } else if (U < 1/2) {
5      X <- 2
6  } else {
7      X <- 3
8  }
9  print(X)
```

Python (NumPy)

```
1  import numpy as np
2
3  U = np.random.uniform(0, 1)
4  if U < 1/6:
5      X = 1
6  elif U < 1/2:
7      X = 2
8  else:
9      X = 3
10
11 print(X)
```

Generation of Random Numbers with Discrete Distributions

Suppose that we want to generate the value of a discrete random variable X having probability mass function

$$P\{X = x_j\} = p_j, \quad j = 0, 1, \dots, \quad \sum_{j=0}^{\infty} p_j = 1.$$

To accomplish this, we generate a random number $U \sim \text{Uniform}(0, 1)$ and set

$$X = \begin{cases} x_0, & \text{if } U < p_0, \\ x_1, & \text{if } p_0 \leq U < p_0 + p_1, \\ \vdots & \end{cases}$$

Since for $0 < a < b < 1$, $P\{a \leq U < b\} = b - a$, we have

$$P\{X = x_j\} = P\left\{\sum_{i=0}^{j-1} p_i \leq U < \sum_{i=0}^j p_i\right\} = p_j$$

and so X has the desired distribution.

Example: If $P\{X = j\} = 1/n$, $j = 0, \dots, n - 1$, then

$$X = j \quad \text{if} \quad \frac{j-1}{n} \leq U < \frac{j}{n}$$

If U is uniformly distributed in $[0, 1]$ we can see that by defining $X = \lfloor nU \rfloor$, the random variable X follows the target distribution.

$$X = \lfloor nU \rfloor$$

where $\lfloor \cdot \rfloor$ denotes the rounding down operation.

Generation of random numbers with discrete distributions

Problem: Consider $n = 50$. Generate a random number from a random variable X with discrete probability distribution given by the p.m.f. $p(i) = 1/n$ for $i = 0, 1, \dots, n - 1$.

Algorithm

1. Generate $U \sim U(0, 1)$
2. Define $X = \lfloor nU \rfloor$
3. Return X

Code

```
1 k <- 50
2 U <- runif(1, min=0, max=1)
3 X <- floor(k*U)
4 print(X)
```

```
1 import numpy as np
2
3 k = 50
4 U = np.random.uniform(0, 1)
5 X = np.floor(k*U)
6 print(int(X))
```

The preceding can be written algorithmically as:

$$\left\{ \begin{array}{l} \text{Generate a random number } U, \\ \text{If } U < p_0, \text{ set } X = x_0 \text{ and stop,} \\ \text{If } U < p_0 + p_1, \text{ set } X = x_1 \text{ and stop,} \\ \text{If } U < p_0 + p_1 + p_2, \text{ set } X = x_2 \text{ and stop,} \\ \vdots \end{array} \right.$$

If the $x_i, i \geq 0$ are ordered so that $x_0 < x_1 < x_2 < \dots$, and if we let F denote the distribution function of X , then $F(x_k) = \sum_{i=0}^k p_i$, and

$$X = x_j \quad \text{if and only if} \quad F(x_{j-1}) \leq U < F(x_j).$$

In other words, after generating U , we determine X by finding the interval $[F(x_{j-1}), F(x_j))$ in which U lies.

This is equivalent to computing the inverse transform $X = F^{-1}(U)$.

Algorithm 1 Generate X with probabilities p_0, \dots, p_n using Uniform(0, 1) numbers

```
1: Input: Probabilities  $p_0, \dots, p_n$ 
2: Draw  $U \sim \text{Uniform}(0, 1)$ 
3: Initialize cumulative probability:  $C \leftarrow 0$ 
4: for  $i = 0$  to  $n$  do
5:   Update cumulative probability:  $C \leftarrow C + p_i$ 
6:   if  $U < C$  then
7:     Set  $X \leftarrow x_i$  and stop
8:   end if
9: end for
10: Output: Sampled value  $X$ 
```

▷ Loop over outcomes

Example: Step-by-Step Calculation

Refer to the previous example: Generate a random number from X with

$$P(X = 1) = 1/6, \quad P(X = 2) = 1/3, \quad P(X = 3) = 1/2$$

Assume we draw $U = 0.42$

- Initialize cumulative probability: $C = 0$
- Step 1: $i = 1, x_1 = 1, p_1 = 1/6 \approx 0.1667$

$$C \leftarrow C + p_1 = 0 + 0.1667 = 0.1667$$

Check: $U < C$? $0.42 < 0.1667 \rightarrow$ **No**

- Step 2: $i = 2, x_2 = 2, p_2 = 1/3 \approx 0.3333$

$$C \leftarrow C + p_2 = 0.1667 + 0.3333 = 0.5$$

Check: $U < C$? $0.42 < 0.5 \rightarrow$ **Yes!** $\rightarrow X = 2$

- Step 3: Not needed, algorithm stops.

Result: $X = 2$

Algorithm 2 Generate X with probabilities p_1, \dots, p_n using cumulative sums

- 1: **Input:** Probabilities p_1, \dots, p_n (with $\sum_{i=1}^n p_i = 1$)
- 2: Sort probabilities in **descending order** and reorder outcomes accordingly:

$$p_{(1)} \geq p_{(2)} \geq \dots \geq p_{(n)}$$

- 3: Compute cumulative sums array:

$$S_j = \sum_{i=1}^j p_{(i)}, \quad j = 1, \dots, n$$

- 4: Draw $U \sim \text{Uniform}(0, 1)$
 - 5: **for** $j = 1$ **to** n **do**
 - 6: **if** $U < S_j$ **then**
 - 7: Set $X \leftarrow x_{(j)}$ and **stop**
 - 8: **end if**
 - 9: **end for**
 - 10: **Output:** Sampled value X
-

What is different here?

By ordering outcomes by probability, the algorithm reduces the **expected number of comparisons**.

The expected number of comparisons is

$$\mathbb{E}[\text{comparisons}] = \sum_{i=1}^n i \cdot p_{(i)}.$$

Example: Geometric Distribution

In some cases however, we would prefer an explicit expression connecting X and U , instead of listing all possible cases.

Let X be a geometric random variable with parameter p if

$$P\{X = i\} = p_i = p^{i-1}(1 - p), \quad i \geq 1$$

From this expression we can see that:

$$\sum_{i=1}^j p_i = \sum_{i=1}^j p^{i-1}(1 - p) = (1 - p) \frac{1 - p^j}{1 - p} = 1 - p^j$$

The event $\left\{ U < \sum_{i=1}^j p_i \right\}$ is equivalent to the event $\left\{ j > \frac{\ln(1 - U)}{\ln(p)} \right\}$

This means that by rounding up $\frac{\ln(1 - U)}{\ln(p)}$ we find the appropriate value of j .

Note that $\log(p) < 0$ for $0 < p < 1$, so the inequality flips as shown above.

Algorithm

1. Generate $U \sim U(0, 1)$
2. Define $X = \left\lceil \frac{\ln(1 - U)}{\ln(p)} \right\rceil$
3. Return X

Remark: The expression $\lceil \cdot \rceil$ denotes the operation of rounding up to the nearest integer.

Algorithm to Generate a Random Number From the Geometric Distribution

Problem: Generate a sample of 10 independent random numbers from the Geometric distribution with parameter $p = 0.4$.

Code

Solution (R)

```
1  n <- 10
2  p <- 0.4
3  U <- runif(n, min=0, max=1)
4  X <- ceiling(log(1-U)/log(p))
5  print(X)
```

Solution (Python)

```
1  import numpy as np
2
3  n = 10
4  p = 0.4
5  U = np.random.uniform(0, 1, n)
6  X = np.ceil(np.log(1-U) / np.log(p))
   .astype(int)
7  print(X)
```

Homework Problems

1. Implement Algorithm 1 and Algorithm 2 and test them on the slide 5 example.
2. Consider a discrete random variable taking values 1, 3, 5 with probabilities $1/9$, $3/9$, and $5/9$ respectively. Write a computer program to generate samples of arbitrary size n from this distribution. Create a histogram with a sample of size 1000.
3. Consider a random variable taking values $1, 3, 5, \dots, 99$ with probabilities $1/50^2, 3/50^2, \dots, 99/50^2$. Write a computer program to generate samples of arbitrary size n from this distribution. Create a histogram with a sample of size 1000.
4. For the slide 5 example, compute the expected number of comparisons using Algorithm 1 and Algorithm 2.
5. Write an algorithm to generate samples from a Poisson distribution:

$$p_i = P\{X = i\} = \frac{e^{-\lambda} \lambda^i}{i!}, \quad i = 0, 1, 2, \dots$$

using the recursive identity:

$$p_{i+1} = \frac{\lambda}{i+1} p_i, \quad i \geq 0$$