

# Chapter 4

Markov Chain Monte Carlo.  
Metropolis Algorithm.

---

Ali Raisolsadat

School of Mathematical and Computational Sciences  
University of Prince Edward Island

## Reminder: Homogeneous Markov Chains

**Definition:** A sequence of random variables  $\{X_t\}_{t \geq 0}$  is a **Markov chain** if

$$P(X_{t+1} = j | X_t = i, X_{t-1}, \dots, X_0) = P(X_{t+1} = j | X_t = i).$$

The process is called **homogeneous** (or time-homogeneous) if the transition probabilities do not depend on  $t$ :

$$P(X_{t+1} = j | X_t = i) = P_{ij}, \quad \text{for all } t.$$

Transition matrix:

$$\mathbf{P} = \begin{pmatrix} P_{11} & P_{12} & \cdots \\ P_{21} & P_{22} & \cdots \\ \vdots & \vdots & \ddots \end{pmatrix}, \quad P_{ij} \geq 0, \quad \sum_j P_{ij} = 1.$$

**n-step transition:**  $P^{(n)} = \mathbf{P}^n$ , with entries  $P_{ij}^{(n)} = P(X_{t+n} = j | X_t = i)$ .

# Stationary Distribution

**Definition:** A probability vector  $\pi$  is a **stationary distribution** of a Markov chain if

$$\pi^T P = \pi^T, \quad \text{and} \quad \sum_i \pi_i = 1, \quad \pi_i \geq 0.$$

**Interpretation:** If  $X_0 \sim \pi$ , then  $X_t \sim \pi$  for all  $t$ . The chain remains in equilibrium under  $\pi$ .

**Convergence:** For many Markov chains,

$$\lim_{t \rightarrow \infty} P(X_t = j | X_0 = i) = \pi_j,$$

independent of the starting state  $i$ .

This limiting behavior holds under certain conditions (irreducibility, aperiodicity, and recurrence).

# Irreducibility, Recurrence, and Ergodicity

**Irreducibility:** A Markov chain is **irreducible** if every state can be reached from every other state:

$$\forall i, j, \exists n \text{ such that } P_{ij}^{(n)} > 0$$

**Recurrence:** State  $i$  is **recurrent** if, starting from  $i$ , the probability of returning to  $i$  is 1:

$$P(\text{return to } i \mid X_0 = i) = 1$$

If the expected return time is finite,  $i$  is **positive recurrent**.

**Ergodicity:** An irreducible, aperiodic, and positive recurrent Markov chain is called **ergodic**. For ergodic chains,

$$\lim_{t \rightarrow \infty} P(X_t = j \mid X_0 = i) = \pi_j$$

and time averages converge to expectations under  $\pi$ :

$$\frac{1}{T} \sum_{t=1}^T f(X_t) \xrightarrow{\text{a.s.}} E_\pi[f(X)]$$

We have been discussing inference using **Markov chains**.

- Concepts such as sampling and stationary distributions are central.
- For **discrete** Markov chains, dynamic programming algorithms can often be used for exact inference (e.g., Hidden Markov Models).

We can also use Markov chains for inference in more general settings.

- The most common framework is **Markov Chain Monte Carlo (MCMC)**.
- MCMC methods are used for **approximate inference**, particularly in complex Bayesian models (e.g., Bayesian logistic regression).

## High-level idea of MCMC:

- We want to compute expectations with respect to a distribution  $p(x)$ , but we cannot generate independent samples directly from  $p$ .
- Construct a **homogeneous Markov chain** whose **stationary distribution** is  $p(x)$ .
- After a suitable burn-in period, use the generated samples  $\{X_t\}_{t=1}^T$  in a Monte Carlo approximation.

## Degenerate Example: “Pointless MCMC”

Consider estimating the expected value of a fair six-sided die. We know analytically that

$$E[X] = \frac{1 + 2 + 3 + 4 + 5 + 6}{6} = 3.5$$

Now suppose we design a “pointless” MCMC algorithm for this trivial problem:

- Start with an initial value  $x_0 \in \{1, 2, 3, 4, 5, 6\}$ , e.g.  $x_0 = 4$ .
- At each iteration  $t$ :
  - Roll the die to propose a new value  $x' \sim \text{Uniform}\{1, 2, 3, 4, 5, 6\}$ .
  - Generate a random number  $u \sim \text{Uniform}[0, 1]$ .
  - If  $u < 0.5$ , **accept**  $x'$ , i.e. set  $x_t = x'$ .
  - Otherwise, **reject**  $x'$  and keep the old value,  $x_t = x_{t-1}$ .

This produces samples from a Markov chain with transition probabilities:

$$q(x_{t-1} \rightarrow x_t) = \frac{1}{2} \mathbf{1}(x_t = x_{t-1}) + \frac{1}{2} \cdot \frac{1}{6} = \begin{cases} \frac{7}{12}, & \text{if } x_t = x_{t-1}, \\ \frac{1}{12}, & \text{if } x_t \neq x_{t-1} \end{cases}$$

# Pointless MCMC in Action

Let's simulate a few steps of the “pointless” MCMC algorithm.

**Setup:** Start with  $x_0 = 4$ .

Step	Proposed Roll	$u$ (Uniform(0,1))	Recorded Value ( $x_t$ )
1	6	0.234	6 (accepted)
2	3	0.612	6 (rejected)
3	2	0.523	6 (rejected)
4	3	0.125	3 (accepted)
5	2	0.433	2 (accepted)

Resulting samples:

$$x_{0:5} = (4, 6, 6, 6, 3, 2, \dots)$$

## Degenerate Example: “Pointless MCMC”

### Observation:

- The chain sometimes repeats values due to rejections.
- Samples are **correlated**, even though the target distribution is uniform.

### Key insight:

- If you run this chain long enough, you will spend roughly  $\frac{1}{6}$  of the time on each outcome.
- The **stationary distribution** is uniform: if we start from a uniform state, either staying there or moving to a uniformly chosen new state keeps the distribution uniform.
- Thus, the stationary distribution of the chain is  $p$  (the uniform distribution).
- The property of constructing a chain whose stationary distribution is  $p$  is the **key idea behind all MCMC methods**.
- It is “pointless” here because we already know how to generate i.i.d. samples from  $p$ . If you can sample directly, you do *not* need MCMC.

# Markov Chain Monte Carlo (MCMC)

**Goal:** Estimate expectations with respect to a distribution  $p(x)$  when direct sampling is difficult.

**Key idea:**

- Construct a **Markov chain** whose **stationary distribution** is  $\pi(x) = p(x)$ .
  - After sufficient iterations (“burn-in”), samples  $x^{(k)}$  approximately follow  $p(x)$ .
  - Notation:  $x^{(1)}$  is the first sampled state,  $x^{(2)}$  the second, ...,  $x^{(n)}$  the  $n$ th.
- Use the dependent Markov chain samples in a Monte Carlo estimator:

$$\mathbb{E}_p[g(X)] \approx \frac{1}{n} \sum_{t=1}^n g(x^{(t)}).$$

- A generalization of the Law of Large Numbers, known as the **Ergodic Theorem**, ensures that as  $n \rightarrow \infty$ :

$$\frac{1}{n} \sum_{t=1}^n g(x^{(t)}) \xrightarrow{\text{a.s.}} \mathbb{E}_p[g(X)].$$

- Convergence is slower than for i.i.d. samples, since the draws are **correlated**.
- The variance of the estimator is typically larger than  $\text{Var}[g(X)]/n$ .
- A widely used method to construct such chains is the **Metropolis–Hastings algorithm**.

## Special Case: Metropolis Algorithm

**Metropolis algorithm:** sampling from a **continuous target distribution**  $p(x)$ . We assume  $p(x)$  can be evaluated up to a normalizing constant:

$$p(x) = \frac{\tilde{p}(x)}{Z}, \quad Z \text{ unknown.}$$

**Algorithm:**

1. Initialize  $x^{(0)}$ .

2. Until we get bored:

2.1 Add zero-mean Gaussian noise to generate proposal  $\hat{x}^{(t)}$

$$\hat{x}^{(t)} = x^{(t-1)} + \epsilon, \quad \epsilon \sim N(0, \sigma^2)$$

2.2 Generate  $u \sim \text{Uniform}[0, 1]$

2.3 Accept the proposal if

$$u \leq \frac{\tilde{p}(\hat{x}^{(t)})}{\tilde{p}(x^{(t-1)})} \frac{\text{probability of proposed}}{\text{probability of current}}$$

and set  $x^{(t)} = \hat{x}^{(t)}$ .

2.4 Otherwise, reject the proposal and set  $x^{(t)} = x^{(t-1)}$

- Proposals that increase the target density are **always accepted**.
- Proposals that decrease the target density **may be accepted or rejected**.
- Under mild conditions (irreducibility, aperiodicity), the chain converges to  $p(x)$ , but convergence may be slow.
- Works even when the normalizing constant  $Z$  is unknown.

# Metropolis Algorithm

---

**Algorithm 1** Metropolis Algorithm for Sampling from  $p(x)$ 

---

**Require:** Initial state  $x^{(0)}$ , number of iterations  $T$ , proposal standard deviation  $\sigma$

**Ensure:** Samples  $\{x^{(t)}\}_{t=1}^T$  approximately distributed according to  $p(x)$

```
1: Set  $x^{(0)}$  as the starting state
2: for  $t = 1$  to  $T$  do
3:   Propose  $\hat{x}^{(t)} = x^{(t-1)} + \epsilon$ , with  $\epsilon \sim N(0, \sigma^2)$ 
4:   Generate  $u \sim \text{Uniform}(0, 1)$ 
5:   if  $u \leq \frac{\tilde{p}(\hat{x}^{(t)})}{\tilde{p}(x^{(t-1)})}$  then
6:     Accept the proposal:  $x^{(t)} \leftarrow \hat{x}^{(t)}$ 
7:   else
8:     Reject the proposal:  $x^{(t)} \leftarrow x^{(t-1)}$ 
9:   end if
10:  end for
```

---

- In practice, we often do not *use all* the samples in our Monte Carlo estimates.
- **Burn-in**: discard the early samples while the chain is far from the stationary distribution.
- **Thinning**: keep only every  $k$ -th sample to reduce autocorrelation between consecutive samples.
- Two common approaches for applying MCMC:
  1. Run a **large number of independent chains** for a short time and use **final states**:
    - Highly parallelizable.
    - Effectively an extreme form of thinning: only one sample per chain is used.
    - Must ensure each chain has reached the stationary distribution (burn-in).
  2. Run a **single chain** for a long time and use **states across time**:
    - Less concern about burn-in if the chain is sufficiently long.
    - Thinning may be needed to reduce autocorrelation among samples.
- Diagnosing whether the chain has reached the stationary distribution can be difficult in practice.

# Homework: Burn-in and Thinning in MCMC

## Problem:

Consider the Metropolis algorithm for sampling from a target distribution

$$p(x) \propto \exp(-0.05x^2) \quad (\text{Normal}(0,10)).$$

1. Simulate a **single Markov chain** of length 200, starting from  $x_0 = -10$ , using a Gaussian random walk proposal with standard deviation 2.
2. Plot the chain over iterations.
3. Highlight the **first 30 iterations** as **burn-in** (shaded in gray).
4. Apply **thinning** by keeping every 5th sample **after burn-in** and mark these samples in red on the plot.
5. In a few sentences, explain why **burn-in** and **thinning** are used in MCMC.
6. **Optional / Bonus: Multiple Chains**
  - Run 1000 independent chains, each starting from a random value in  $[-10,10]$ .
  - Apply burn-in and thinning as before.
  - Combine all thinned samples from all chains and plot a histogram.
  - Comment on how the histogram approximates the target distribution.