

Chapter 4

Markov Chain Monte Carlo.
Metropolis Algorithm.

Ali Raisolsadat

School of Mathematical and Computational Sciences
University of Prince Edward Island

We have been discussing inference using **Markov chains**.

- Concepts such as sampling and stationary distributions are central.
- For **discrete** Markov chains, dynamic programming algorithms can often be used for exact inference (e.g., Hidden Markov Models).

We can also use Markov chains for inference in more general settings.

- The most common framework is **Markov Chain Monte Carlo (MCMC)**.
- MCMC methods are used for **approximate inference**, particularly in complex Bayesian models (e.g., Bayesian logistic regression).

High-level idea of MCMC:

- We want to compute expectations with respect to a distribution $p(x)$, but we cannot generate independent samples directly from p .
- Construct a **homogeneous Markov chain** whose **stationary distribution** is $p(x)$.
- After a suitable burn-in period, use the generated samples $\{X_t\}_{t=1}^T$ in a Monte Carlo approximation.

Definition: A sequence of random variables $\{X_t\}_{t \geq 0}$ is a **Markov chain** if

$$P(X_{t+1} = j | X_t = i, X_{t-1}, \dots, X_0) = P(X_{t+1} = j | X_t = i).$$

The process is called **homogeneous** (or time-homogeneous) if the transition probabilities do not depend on t :

$$P(X_{t+1} = j | X_t = i) = P_{ij}, \quad \text{for all } t.$$

Transition matrix:

$$\mathbf{P} = \begin{pmatrix} P_{11} & P_{12} & \cdots \\ P_{21} & P_{22} & \cdots \\ \vdots & \vdots & \ddots \end{pmatrix}, \quad P_{ij} \geq 0, \quad \sum_j P_{ij} = 1.$$

n-step transition: $P^{(n)} = \mathbf{P}^n$, with entries $P_{ij}^{(n)} = P(X_{t+n} = j | X_t = i)$.

Definition: A probability vector π is a **stationary distribution** of a Markov chain if

$$\pi^\top \mathbf{P} = \pi^\top, \quad \sum_i \pi_i = 1, \quad \pi_i \geq 0.$$

Interpretation: If $X_0 \sim \pi$ then $X_t \sim \pi$ for all t . The chain remains in equilibrium under π .

Convergence: For many Markov chains

$$\lim_{t \rightarrow \infty} P(X_t = j \mid X_0 = i) = \pi_j$$

Irreducibility, Recurrence, and Ergodicity

Irreducibility: A Markov chain is **irreducible** if every state can be reached from every other state:

$$\forall i, j, \exists n \text{ such that } P_{ij}^{(n)} > 0$$

Recurrence: State i is **recurrent** if, starting from i , the probability of returning to i is 1:

$$P(\text{return to } i \mid X_0 = i) = 1$$

If the expected return time is finite, i is **positive recurrent**.

Ergodicity: An irreducible, aperiodic (gcd is 1 for all possible return times to a state), and positive recurrent Markov chain is called **ergodic**. For ergodic chains,

$$\lim_{t \rightarrow \infty} P(X_t = j \mid X_0 = i) = \pi_j$$

and time averages converge to expectations under π :

$$\frac{1}{T} \sum_{t=1}^T f(X_t) \xrightarrow{\text{a.s.}} E_\pi[f(X)]$$

Degenerate Example: “Pointless MCMC”

Consider estimating the expected value of a fair six-sided die. We know analytically that

$$E[X] = \frac{1 + 2 + 3 + 4 + 5 + 6}{6} = 3.5$$

Now suppose we design a “pointless” MCMC algorithm for this trivial problem:

- Start with an initial value $x_0 \in \{1, 2, 3, 4, 5, 6\}$, e.g. $x_0 = 4$.
- At each iteration t :
 - Roll the die to propose a new value $x' \sim \text{Uniform}\{1, 2, 3, 4, 5, 6\}$.
 - Generate a random number $u \sim \text{Uniform}[0, 1]$.
 - If $u < 0.5$, **accept** x' , i.e. set $x_t = x'$.
 - Otherwise, **reject** x' and keep the old value, $x_t = x_{t-1}$.

This produces samples from a Markov chain with transition probabilities:

$$q(x_{t-1} \rightarrow x_t) = \frac{1}{2} \mathbf{1}(x_t = x_{t-1}) + \frac{1}{2} \cdot \frac{1}{6} = \begin{cases} \frac{7}{12}, & \text{if } x_t = x_{t-1}, \\ \frac{1}{12}, & \text{if } x_t \neq x_{t-1} \end{cases}$$

Let's simulate a few steps of the “pointless” MCMC algorithm.


Setup: Start with $x_0 = 4$.

Step	Proposed Roll	u (Uniform(0,1))	Recorded Value (x_t)
1	6	0.234	6 (accepted)
2	3	0.612	6 (rejected)
3	2	0.523	6 (rejected)
4	3	0.125	3 (accepted)
5	2	0.433	2 (accepted)

Resulting samples:

$$x_{0:5} = (4, 6, 6, 6, 3, 2, \dots)$$

Convergence to the Actual Expectation - Pointless MCMC in Action



Degenerate Example: “Pointless MCMC”

Observation:

- The chain sometimes repeats values due to rejections.
- Samples are **correlated**, even though the target distribution is uniform.

Key insight:

- If you run this chain long enough, you will spend roughly $\frac{1}{6}$ of the time on each outcome.
- The **stationary distribution** is uniform: if we start from a uniform state, either staying there or moving to a uniformly chosen new state keeps the distribution uniform.
- Thus, the stationary distribution of the chain is p (the uniform distribution).
- The property of constructing a chain whose stationary distribution is p is the **key idea behind all MCMC methods**.
- It is “pointless” here because we already know how to generate i.i.d. samples from p . If you can sample directly, you do *not* need MCMC.

Markov Chain Monte Carlo (MCMC)

Goal: Estimate expectations with respect to a distribution $p(x)$ when direct sampling is difficult.

Key idea:

- Construct a **Markov chain** whose **stationary distribution** is $\pi(x) = p(x)$.
 - After sufficient iterations ("burn-in"), samples $x^{(k)}$ approximately follow $p(x)$.
 - Notation: $x^{(1)}$ is the first sampled state, $x^{(2)}$ the second, \dots , $x^{(n)}$ the n th.
- Use the dependent Markov chain samples in a Monte Carlo estimator:

$$\mathbb{E}_p[g(X)] \approx \frac{1}{n} \sum_{t=1}^n g(x^{(t)}).$$

- A generalization of the Law of Large Numbers, known as the **Ergodic Theorem**, ensures that as $n \rightarrow \infty$:

$$\frac{1}{n} \sum_{t=1}^n g(x^{(t)}) \xrightarrow{\text{a.s.}} \mathbb{E}_p[g(X)].$$

- Convergence is slower than for i.i.d. samples, since the draws are **correlated**.
- The variance of the estimator is typically larger than $\text{Var}[g(X)]/n$.
- A widely used method to construct such chains is the **Metropolis–Hastings algorithm**.

Special Case: Metropolis Algorithm

Metropolis algorithm: Sampling from a **continuous target distribution** $p(x)$. We assume $p(x)$ can be evaluated up to a normalizing constant:

$$p(x) = \frac{\tilde{p}(x)}{Z}, \quad Z \text{ unknown.}$$

Algorithm:

1. Initialize $x^{(0)}$.

2. Until we get bored:

2.1 Add zero-mean Gaussian noise to generate proposal $\hat{x}^{(t)}$

$$\hat{x}^{(t)} = x^{(t-1)} + \epsilon, \quad \epsilon \sim N(0, \sigma^2)$$

2.2 Generate $u \sim \text{Uniform}[0, 1]$

2.3 Accept the proposal if

$$u \leq \frac{\tilde{p}(\hat{x}^{(t)})}{\tilde{p}(x^{(t-1)})}, \quad u \leq \frac{\text{probability of proposed}}{\text{probability of current}}$$

and set $x^{(t)} = \hat{x}^{(t)}$.

2.4 Otherwise, reject the proposal and set $x^{(t)} = x^{(t-1)}$

- Proposals that increase the target density are **always accepted**.
- Proposals that decrease the target density **may be accepted or rejected**.
- Under mild conditions (irreducibility, aperiodicity), the chain converges to $p(x)$, but convergence may be slow.
- Works even when the normalizing constant Z is unknown.

1D Graphical Intuition: Why Some Proposals Are Accepted

Consider a one-dimensional target density $p(x)$.

- Suppose the chain is currently at $x^{(t-1)}$.
- A proposal \hat{x} is generated by adding Gaussian noise:

$$\hat{x} = x^{(t-1)} + \epsilon \quad \epsilon \sim N(0, \sigma^2)$$

Two possible situations

- **Uphill move:** $p(\hat{x}) \geq p(x^{(t-1)})$

$$\frac{\tilde{p}(\hat{x})}{\tilde{p}(x^{(t-1)})} \geq 1$$

The proposal is accepted with probability 1.

- **Downhill move:** $p(\hat{x}) < p(x^{(t-1)})$

$$\frac{\tilde{p}(\hat{x})}{\tilde{p}(x^{(t-1)})} < 1$$

The proposal is accepted with probability

$$\alpha = \frac{\tilde{p}(\hat{x})}{\tilde{p}(x^{(t-1)})}$$

Metropolis Algorithm

Algorithm 1 Metropolis Algorithm for Sampling from $p(x)$


Require: Initial state $x^{(0)}$, number of iterations T , proposal standard deviation σ

Ensure: Samples $\{x^{(t)}\}_{t=1}^T$ approximately distributed according to $p(x)$

```
1: Set  $x^{(0)}$  as the starting state
2: for  $t = 1$  to  $T$  do
3:   Propose  $\hat{x}^{(t)} = x^{(t-1)} + \epsilon$ , with  $\epsilon \sim N(0, \sigma^2)$ 
4:   Generate  $u \sim \text{Uniform}(0, 1)$ 
5:   if  $u \leq \frac{\tilde{p}(\hat{x}^{(t)})}{\tilde{p}(x^{(t-1)})}$  then
6:     Accept the proposal:  $x^{(t)} \leftarrow \hat{x}^{(t)}$ 
7:   else
8:     Reject the proposal:  $x^{(t)} \leftarrow x^{(t-1)}$ 
9:   end if
10: end for
```


Metropolis Algorithm - Convergence from Uniform to Target Distribution

From top-left to bottom-right: histograms of 1000 independent Markov chains
Target: $\text{Normal}(0,10)$; Proposal: $\text{Normal}(0,5)$




Metropolis Algorithm 2D

$M = (-0.206, 0.490)$; $N_p = 1000$, Acceptance Rate = $N_a/N_p = 0.69$



- In practice, we often do not *use all* the samples in our Monte Carlo estimates.
- **Burn-in:** discard the early samples while the chain is far from the stationary distribution.
- **Thinning:** keep only every k -th sample to reduce autocorrelation between consecutive samples.
- Two common approaches for applying MCMC:
 1. Run a **large number of independent chains** for a short time and use **final states**:
 - Highly parallelizable.
 - Effectively an extreme form of thinning: only one sample per chain is used.
 - Must ensure each chain has reached the stationary distribution (burn-in).
 2. Run a **single chain** for a long time and use **states across time**:
 - Less concern about burn-in if the chain is sufficiently long.
 - Thinning may be needed to reduce autocorrelation among samples.
- Diagnosing whether the chain has reached the stationary distribution can be difficult in practice.

Visualization of Burn-in and Thinning in MCMC



Problem: Consider the Metropolis algorithm for sampling from a target distribution

$$p(x) \propto \exp(-0.05x^2) \quad (\text{Normal}(0,10)).$$

1. Simulate a **single Markov chain** of length 200, starting from $x_0 = -10$, using a Gaussian random walk proposal with standard deviation 2.
2. Plot the chain over iterations.
3. Highlight the **first 30 iterations** as **burn-in** (shaded in gray).
4. Apply **thinning** by keeping every 5th sample **after burn-in** and mark these samples in red on the plot.
5. In a few sentences, explain why **burn-in** and **thinning** are used in MCMC.
6. **Optional / Bonus: Multiple Chains**
 - Run 1000 independent chains, each starting from a random value in $[-10,10]$.
 - Apply burn-in and thinning as before.
 - Combine all thinned samples from all chains and plot a histogram.
 - Comment on how the histogram approximates the target distribution.